



澳門大學
UNIVERSIDADE DE MACAU
UNIVERSITY OF MACAU

Learning to Reach Goals via Diffusion

GongYing

gong.ying@connect.umac.mo

2024.1.5

Contents

1. Introduction
2. Preliminaries
3. Reaching Goals via Diffusion
4. Goal-conditioned Diffusion Policy
5. Experiments



1. Introduction

- In RL, agents learn behaviors supervised by only a **reward** function.
- Goal-conditioned RL (GCRL) aims to learn **general** policies that can reach arbitrary target states or goals within an environment requiring **no extensive** retraining. But its disadvantage is **sparse** rewards.
- Thus combining GCRL and offline RL is good for generalization and data efficiency.

Challenges:

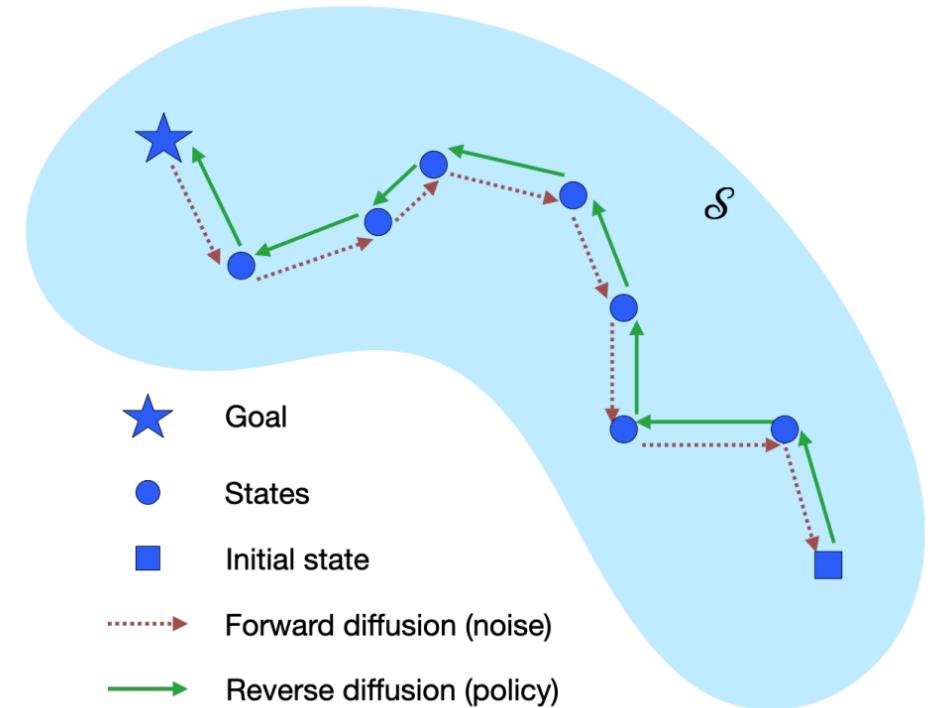
- However, in **offline** RL many methods rely on value function, which estimates the expected discounted return associated with a given **state-action** pair. During training, policies generate actions **not in** the offline dataset (for without interaction), leading to inaccuracies and even diverging policies.
- To solve the problem of limited state-goal pairs, **hindsight** relabeling is employed. But it only generate state-goal pairs within the same trajectory, resulting in over-fitting.



1. Introduction

To solve the problems mentioned, the paper draws inspiration from **diffusion** models.

They construct trajectories that move **away** (noise) from desired goals during the learning process. Then a policy is trained to **reverse** (denoise) the trajectories. Thus the policy learns to reach any predefined goal state to from arbitrary initial states.




2. Preliminaries

I. Diffusion probabilistic models

- The forward diffusion process is to add Gaussian noise to the real image X_0 at each timestep, leading to the final noised image X_T . At timestep t , a constant β_t is given, we have:

$$q(X_t|X_{t-1}) = N(X_t; \sqrt{1 - \beta_t}X_{t-1}, \beta_t I),$$
$$q(X_{1:T}|X_0) = \prod_{t=1}^T q(X_t|X_{t-1}).$$



2. Preliminaries


I. Diffusion probabilistic models

- Then a denoising function is learned to reverse the forward diffusion process.

$$p_{\theta}(X_{t-1}|X_t) = N(X_{t-1}; \mu_{\theta}(X_t, t), \Sigma_{\theta}(X_t, t)),$$

$$p_{\theta}(X_{0:T}|X_0) = p(X_T) \prod_{t=1}^T p_{\theta}(X_{t-1}|X_t),$$

where μ_{θ} and Σ_{θ} can be neural networks.



2. Preliminaries

II. Goal-conditioned RL

- The RL problem can be described using a Markov Decision Process (MDP), denoted by $(S, A, P, r, \mu, \gamma)$, where $\mu(s)$ is the initial state distribution.
- While Goal-conditioned RL additionally considers a goal space $G = \{\phi(s) | s \in S\}$ where $\phi: S \rightarrow G$ is a known state-to-goal mapping. Now reward function depends on the goal and can be sparse and **binary** defined as $r(s, a, g) = \mathbb{1}[||\phi(s) - g||_2^2 \leq \delta]$, where δ is some threshold distance.
- A goal-conditioned policy is denoted by $\pi: S \times G \rightarrow A$, and given a distribution over desired goals $p(g)$, an optimal policy π^* aims to maximize the expected return:

$$J(\pi) = \mathbb{E}_{g \sim p(g), s_0 \sim \mu(s_0), a_t \sim \pi(\cdot | s_t, g), s_{t+1} \sim P(\cdot | s_t, a_t)} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t, g)].$$



3. Reaching Goals via Diffusion

Consider a goal-augmented MDP $(S, A, G, P, r, \mu, \gamma)$ with goals $g \in G$ sampled from unknown goal distribution $g \sim p(g)$. Goal-conditioned RL aims to learn a policy that can learn an optimal path from any state $s \in S$ to the desired goal g .

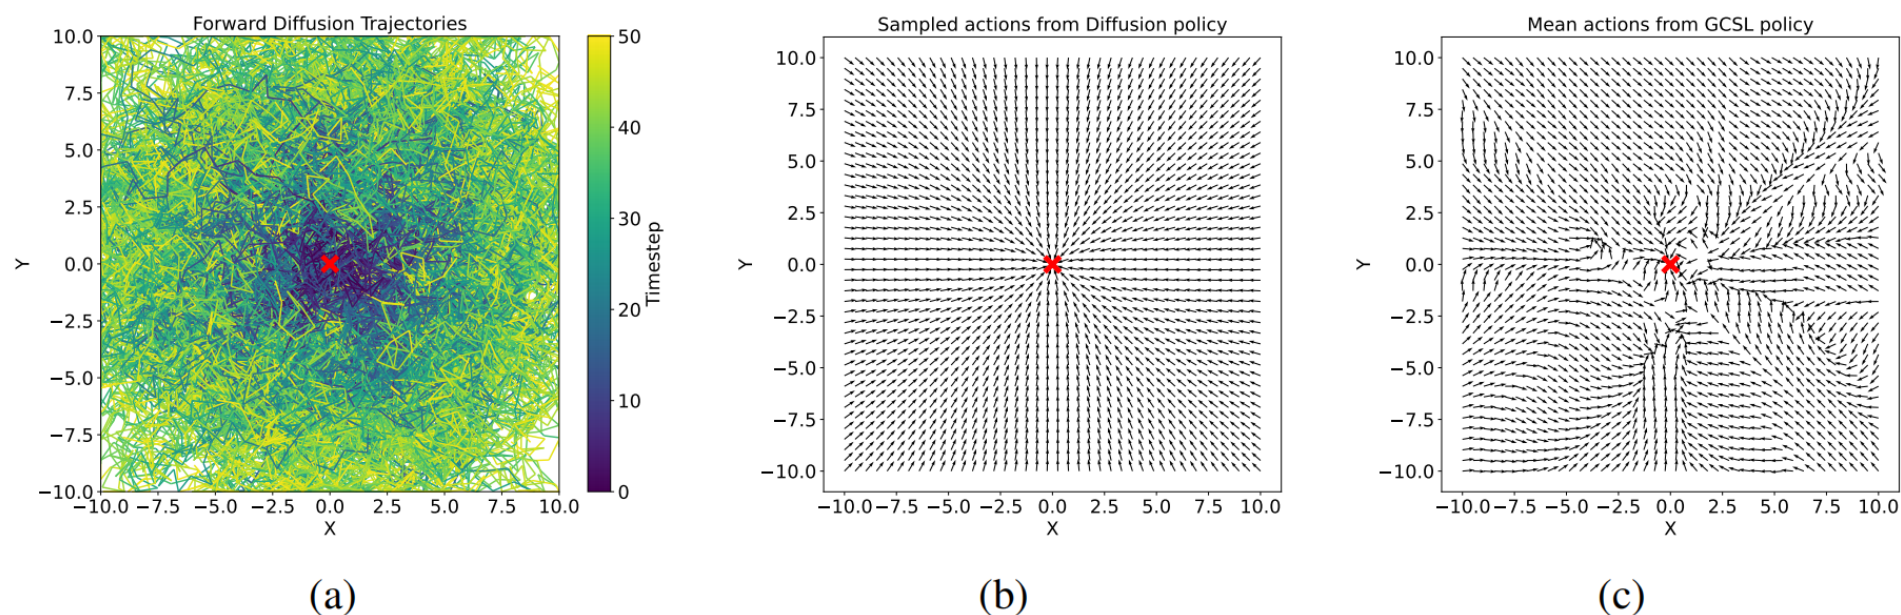


Figure 2: (a) Visualization of trajectories starting from the goal **X** generated during the forward process, (b) Predicted actions from policy trained via diffusion, (c) Predicted actions from policy trained using GCSL.

3. Reaching Goals via Diffusion

During training, the **time horizon** indicates the time difference between the current and desired goal states.

For $h = 1$, the policy always takes the most **direct** path to the goal regardless of the input state. For larger values of the time horizon, the policy has a high variance close to the goal and a low variance for the optimal action further away.

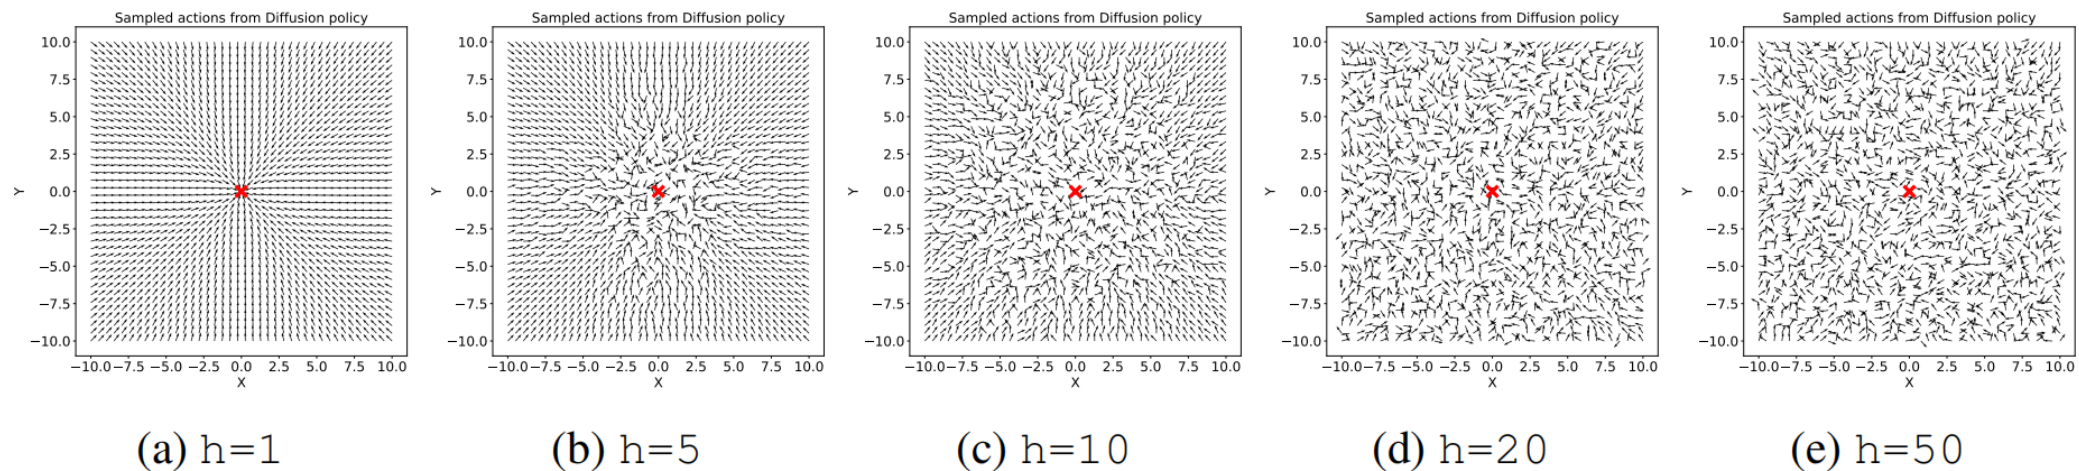


Figure 3: Actions sampled from the trained policy, showcasing the effect of time horizon during evaluation.

4. Goal-conditioned Diffusion Policy

Nearest-neighbor trajectory stitching

The forward diffusion constructs trajectories walking away from the goal to provide training data for the policy. In order for this strategy to be effective, we want to generate **as many state-goal pairs as possible** to help the policy generalize well. Hindsight relabeling can generate positive goal-conditioned observations by replacing the desired goals with achieved goals (e.g. A robot is supposed to move a ball from A to B, but it moved it to C. Then we can assume the task is from A to C, leading to a successful task).

Algorithm 1 Nearest-neighbor Trajectory Stitching

Input: Dataset \mathcal{D} , distance threshold δ , number of new trajectories to collect M
Output: Augmented dataset \mathcal{D}_{new}
 $\mathcal{D}_{\text{new}} \leftarrow \mathcal{D}$
Construct ball tree T for all states
for $m \leftarrow 1, \dots, M$ **do**
 Sample random final state s_T from \mathcal{D}
 $\tau_{\text{new}} \leftarrow \{s_T\}$
 $s_{\text{current}} \leftarrow s_T$
 for $t \leftarrow T, \dots, 1$ **do**
 $s_{\text{neighbor}}, d \leftarrow T.\text{query}(s_{\text{current}}, k = 1)$
 if $d \leq \delta$ **then**
 Add preceding $(s_{\text{prev}}, a_{\text{prev}})$ from s_{neighbor} to τ_{new}
 else
 Add preceding $(s_{\text{prev}}, a_{\text{prev}})$ from s_{current} to τ_{new}
 end if
 $s_{\text{current}} \leftarrow s_{\text{prev}}$
 end for
 $\mathcal{D}_{\text{new}} \leftarrow \mathcal{D}_{\text{new}} \cup \tau_{\text{new}}$
end for
Return: \mathcal{D}_{new}




Figure 5: Trajectory stitching.

5. Experiments

Table 1: Discounted returns, averaged over 10 seeds.

Task Name		Ours			Offline GCRL				Diffusion-based	
		Merlin	Merlin-P	Merlin-NP	GoFAR	WGCSL	GCSL	AM	DD	g-DQL
Expert	PointReach	29.26±0.04	29.17±0.15	29.30±0.05	27.18±0.65	25.91±0.87	22.85±1.26	26.14±1.11	10.03±0.88	28.65±0.44
	PointRooms	25.38±0.37	25.25±0.07	25.42±0.32	20.40±1.00	19.90±0.99	18.28±2.29	23.24±1.58	5.84±2.67	27.53±0.57
	Reacher	22.75±0.59	23.25±0.17	24.97±0.54	22.51±0.82	23.35±0.64	20.05±1.37	22.36±1.03	4.39±1.08	22.54±1.42
	SawyerReach	26.89±0.07	25.05±0.60	27.35±0.06	22.82±1.15	22.07±1.46	19.20±1.79	23.56±0.33	3.39±0.75	24.17±0.01
	SawyerDoor	26.18±2.19	25.75±0.97	26.15±2.08	23.62±0.35	23.92±1.10	20.12±1.33	26.39±0.42	7.85±0.77	24.81±0.38
	FetchReach	30.29±0.03	30.26±0.02	30.42±0.04	29.21±0.26	28.17±0.38	23.68±1.07	29.08±0.12	1.55±0.68	28.71±0.15
	FetchPush	19.91±1.20	2.23±2.20	21.58±1.63	22.41±1.69	22.22±1.51	17.58±1.47	19.86±3.16	5.49±2.85	17.82±0.55
	FetchPick	19.66±0.78	1.43±1.01	20.41±0.92	19.79±1.12	18.32±1.56	12.95±1.90	17.04±3.81	2.76±0.64	14.45±0.61
	FetchSlide	4.19±1.89	0.00±0.00	4.95±2.02	3.34±1.01	5.17±3.17	1.67±1.41	3.31±1.46	1.21±0.59	0.98±0.59
	HandReach	22.11±0.55	0.00±0.00	24.93±0.49	15.39±6.37	18.05±5.12	0.15±0.11	0.00±0.00	0.00±0.00	0.00±0.00
Average Rank		2.7	5.3	1.6	4.5	4.6	7.3	5.0	8.3	5.1
Random	PointReach	29.26±0.04	29.21±0.08	29.31±0.04	23.96±0.93	25.76±0.96	17.74±1.84	25.55±0.57	10.12±0.72	22.65±1.57
	PointRooms	24.80±0.36	24.07±0.19	25.16±0.59	18.09±4.13	19.41±1.01	14.69±2.51	19.10±1.39	5.76±2.99	20.88±0.96
	Reacher	21.09±0.65	16.65±0.48	22.24±0.54	25.10±0.68	22.98±0.91	10.62±2.30	23.70±0.62	4.74±0.36	6.06±0.84
	SawyerReach	26.70±0.14	25.46±0.12	26.86±0.07	19.48±1.39	21.32±1.40	8.78±2.59	25.29±0.35	3.46±0.86	2.84±0.05
	SawyerDoor	19.05±0.66	18.26±1.18	21.69±2.36	20.69±2.14	19.58±3.55	12.47±3.08	10.82±1.67	7.92±0.86	14.77±0.51
	FetchReach	30.42±0.04	30.38±0.02	30.42±0.04	28.34±0.98	27.94±0.30	18.96±1.77	27.11±0.22	1.71±0.77	1.21±0.46
	FetchPush	5.21±0.43	5.08±0.32	7.22±0.35	6.99±1.27	5.35±3.36	4.22±2.19	4.53±1.94	4.49±1.34	5.35±0.23
	FetchPick	3.75±0.18	3.02±0.16	4.36±0.19	3.81±3.71	1.87±1.59	0.81±0.82	3.08±1.35	2.16±0.75	2.17±0.18
	FetchSlide	2.67±0.35	0.00±0.00	3.15±0.14	1.32±1.22	1.04±0.98	0.24±0.27	1.12±0.39	1.31±0.52	0.00±0.00
	HandReach	14.89±2.54	0.00±0.00	17.61±3.06	0.08±0.07	2.54±1.42	1.41±0.51	0.00±0.00	0.00±0.00	0.00±0.00
Average Rank		2.8	4.8	1.3	3.8	4.5	7.3	5.3	7.7	6.7

Maximum trajectory length=50, reward=1 if goal is reached (else 0).

Merlin-P uses a proposed learned parametric reverse dynamics model and reverse policy, to generate additional diffusion trajectories. Merlin-NP uses the trajectory stitching method.



5. Experiments




Figure 6: Discounted returns for each dataset with different values of (a) hindsight ratio and (b) time horizon during evaluation. Values are normalized with respect to the maximum value in each row.

Thank you.

